🏠   Modules   Data connection   Retrievers   MultiQueryRetriever

MultiQueryRetriever

Distance-based vector database retrieval embeds (represents) queries in high-dimensional space and finds similar embedded documents based on "distance". But, retrieval may produce difference results with subtle changes in query wording or if the embeddings do not capture the semantics of the data well. Prompt engineering / tuning is sometimes done to manually address these problems, but can be tedious.

The `MultiQueryRetriever` automates the process of prompt tuning by using an LLM to generate multiple queries from different perspectives for a given user input query. For each query, it retrieves a set of relevant documents and takes the unique union across all queries to get a larger set of potentially relevant documents. By generating multiple perspectives on the same question, the `MultiQueryRetriever` might be able to overcome some of the limitations of the distance-based retrieval and get a richer set of results.

```python
# Build a sample vectorDB
from langchain.vectorstores import Chroma
from langchain.document_loaders import WebBaseLoader
from langchain.embeddings.openai import OpenAIEmbeddings
from langchain.text_splitter import RecursiveCharacterTextSplitter

# Load blog post
loader = WebBaseLoader("https://lilianweng.github.io/posts/2023-06-23-agent/")
data = loader.load()

# Split
text_splitter = RecursiveCharacterTextSplitter(chunk_size=500, chunk_overlap=0)
splits = text_splitter.split_documents(data)

# VectorDB
embedding = OpenAIEmbeddings()
vectordb = Chroma.from_documents(documents=splits, embedding=embedding)
```

API Reference:

- Chroma from `langchain.vectorstores`

- WebBaseLoader from `langchain.document_loaders`
- OpenAIEmbeddings from `langchain.embeddings.openai`
- RecursiveCharacterTextSplitter from `langchain.text_splitter`

`Simple usage`

Specify the LLM to use for query generation, and the retriver will do the rest.

```python
from langchain.chat_models import ChatOpenAI
from langchain.retrievers.multi_query import MultiQueryRetriever

question = "What are the approaches to Task Decomposition?"
llm = ChatOpenAI(temperature=0)
retriever_from_llm = MultiQueryRetriever.from_llm(
    retriever=vectordb.as_retriever(), llm=llm
)
```

API Reference:

- ChatOpenAI from `langchain.chat_models`
- MultiQueryRetriever from `langchain.retrievers.multi_query`

```python
# Set logging for the queries
import logging

logging.basicConfig()
logging.getLogger("langchain.retrievers.multi_query").setLevel(logging.INFO)
```

```python
unique_docs = retriever_from_llm.get_relevant_documents(query=question)
len(unique_docs)
```

```
    INFO:langchain.retrievers.multi_query:Generated queries: ['1. How can
Task Decomposition be approached?', '2. What are the different methods for
Task Decomposition?', '3. What are the various approaches to decomposing
tasks?']
```

5

`Supplying your own prompt`

You can also supply a prompt along with an output parser to split the results into a list of queries.

```python
from typing import List
from langchain import LLMChain
from pydantic import BaseModel, Field
from langchain.prompts import PromptTemplate
from langchain.output_parsers import PydanticOutputParser


# Output parser will split the LLM result into a list of queries
class LineList(BaseModel):
    # "lines" is the key (attribute name) of the parsed output
    lines: List[str] = Field(description="Lines of text")


class LineListOutputParser(PydanticOutputParser):
    def __init__(self) -> None:
        super().__init__(pydantic_object=LineList)

    def parse(self, text: str) -> LineList:
        lines = text.strip().split("\n")
        return LineList(lines=lines)


output_parser = LineListOutputParser()

QUERY_PROMPT = PromptTemplate(
    input_variables=["question"],
    template="""You are an AI language model assistant. Your task is to
generate five
    different versions of the given user question to retrieve relevant
documents from a vector
    database. By generating multiple perspectives on the user question,
your goal is to help
    the user overcome some of the limitations of the distance-based
```

```python
    similarity search.
        Provide these alternative questions seperated by newlines.
        Original question: {question}""",
)
llm = ChatOpenAI(temperature=0)

# Chain
llm_chain = LLMChain(llm=llm, prompt=QUERY_PROMPT,
output_parser=output_parser)

# Other inputs
question = "What are the approaches to Task Decomposition?"
```

API Reference:

- PromptTemplate from `langchain.prompts`

- PydanticOutputParser from `langchain.output_parsers`

```python
# Run
retriever = MultiQueryRetriever(
    retriever=vectordb.as_retriever(), llm_chain=llm_chain,
parser_key="lines"
)  # "lines" is the key (attribute name) of the parsed output

# Results
unique_docs = retriever.get_relevant_documents(
    query="What does the course say about regression?"
)
len(unique_docs)
```

```
    INFO:langchain.retrievers.multi_query:Generated queries: ["1. What is
the course's perspective on regression?", '2. Can you provide information
on regression as discussed in the course?', '3. How does the course cover
the topic of regression?', "4. What are the course's teachings on
regression?", '5. In relation to the course, what is mentioned about
regression?']
```

11